

Comunicações Industriais

Industrial Communications

2022/2023

Support information for lab assignment 1a

TCP sockets in C

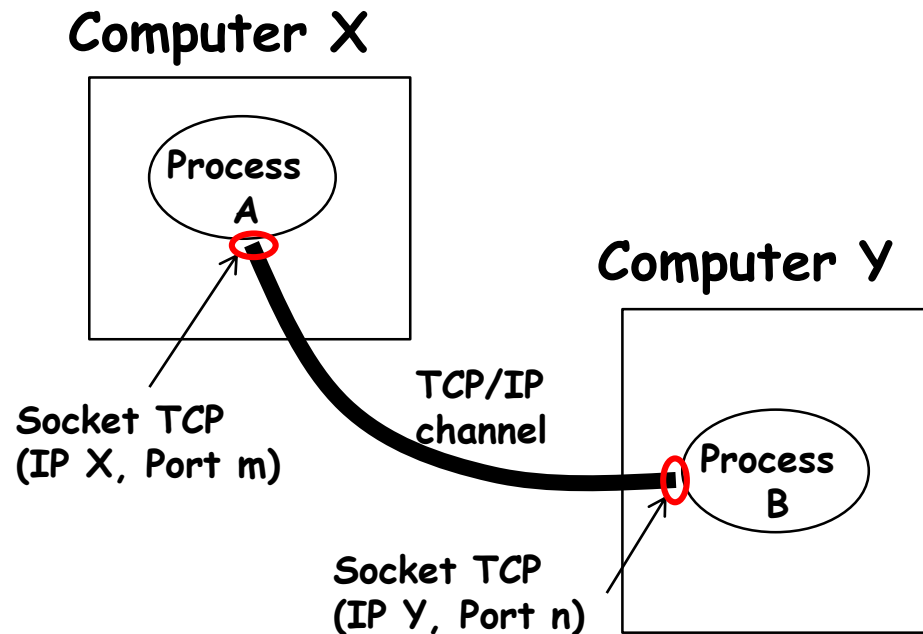
Learning objectives

- Understand the concept of communication channel.
- Understand the construction of channels with *sockets*.
- Ability to create TCP sockets in C for IP networks.

Communication channels and sockets

- **Sockets**

- technology to create **logical communication channels** between **processes** (normally in different computers)
- Sockets = channel end-points



Communication channels and sockets

- **Sockets**
 - Exist for multiple protocol families, protocols and channel types
- We will use just
 - **IP** (Internet) protocol family and the **TCP** protocol
 - Programming language: **C** environment: **Linux** (or Cygwin)
 - » Libraries `<sys/socket.h>` `<netinet/in.h>` `<sys/types.h>`
 - » Online manual pages, many resources on the web...

Creating a TCP channel

• TCP channels

- Follow a **connection-oriented** model

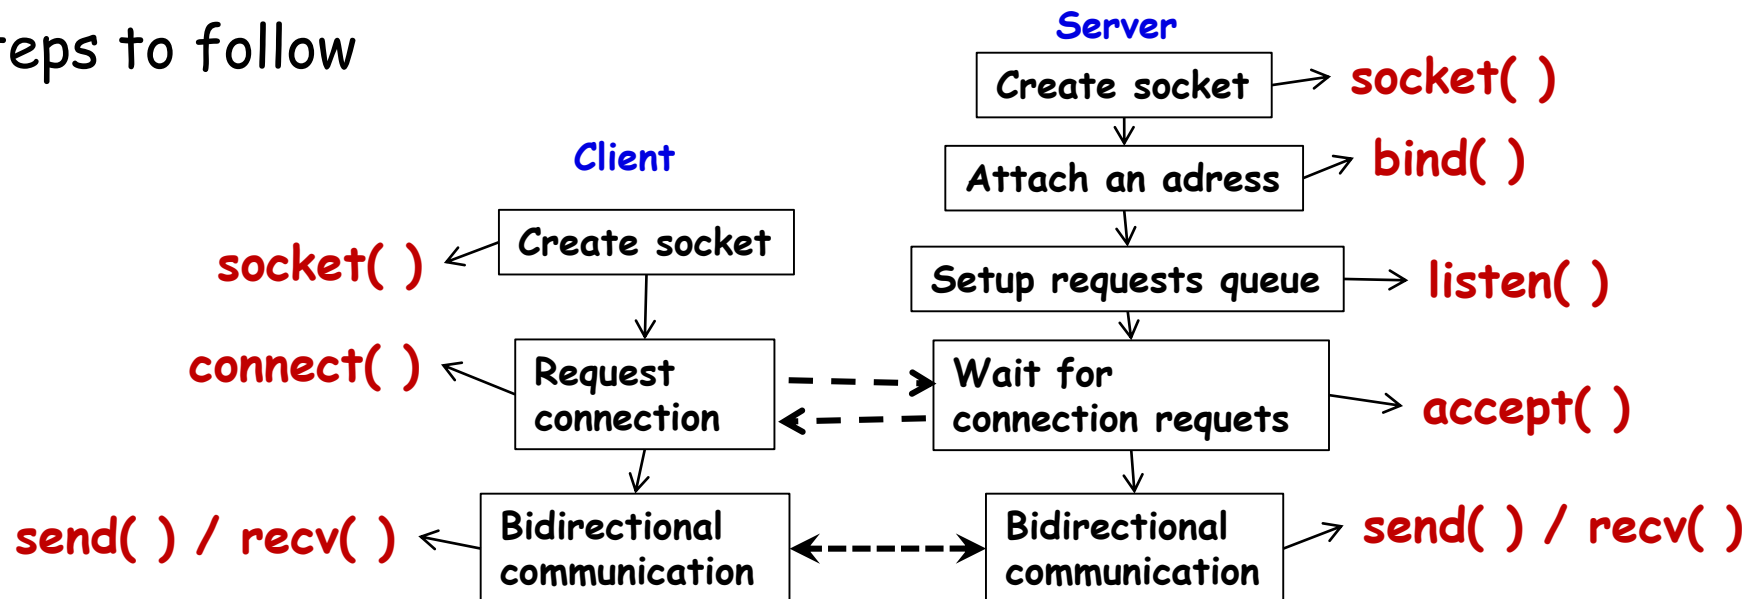
- » First set up a **connection** and **only after** data can be **communicated**

- » Connections are **asymmetric**

Server: passive, waits for connection requests

Client: active, takes the initiative to request connections

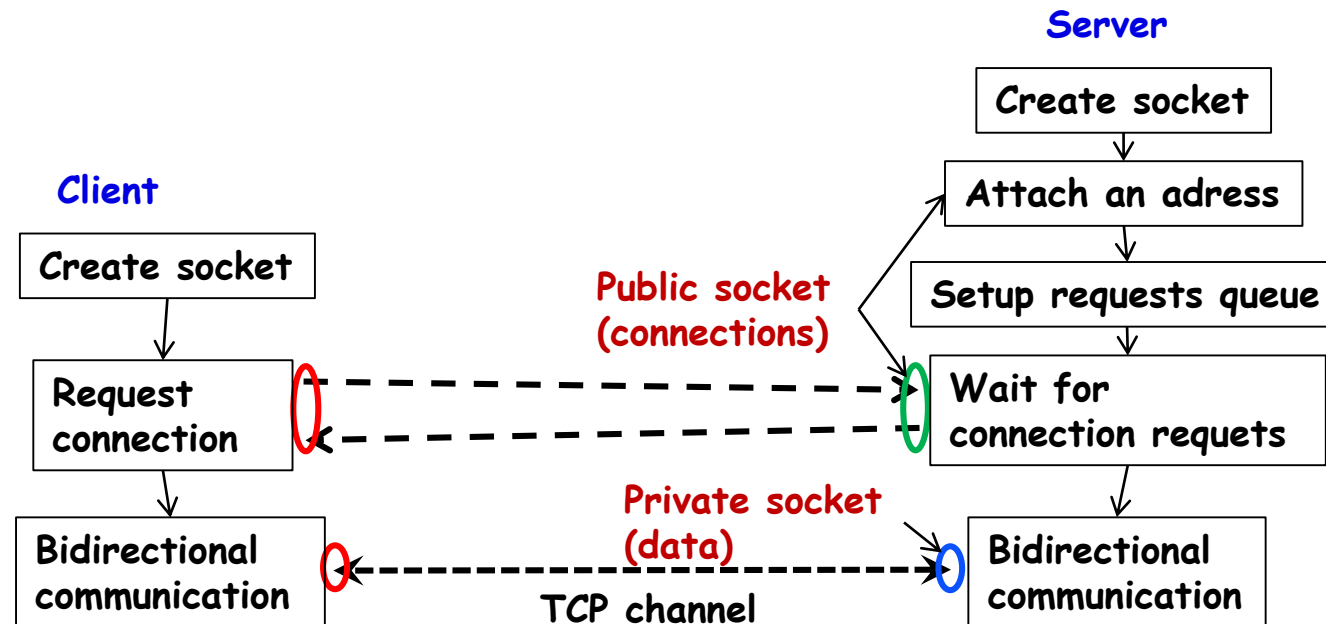
• Steps to follow



Channels and TCP sockets

• TCP channels

- Identified by the (IP Address, TCP Port) pairs of their end-points
 - » **Server:** socket to wait for **connections** + socket for **data communication**
 - » **Client:** uses same socket



Creating a socket

- `int socket (int domain, int type, int protocol)`
 - **domain**: protocol family"- in practice **PF_INET** for the Internet;
 - **type**: channel properties - in practice **SOCK_STREAM** for TCP;
 - **protocol**: which protocol - in practice 0 = **IPPROTO_TCP**
- returns a **socket local identifier**

Used in the server and the client

Attaching an address to a socket

- `int bind (int sockfd, struct sockaddr *my_addr, socklen_t addrlen)`
 - **sockfd**: socket local identifier (returned by *socket()*)
 - **my_addr**: pointer to structure with address to attach
 - **addrlen**: length of the structure pointed to by *my_addr*

Mandatory in the server (to publicize the server)

Optional in the client (normally not used)

Defining addresses and data types

- `struct sockaddr_in { /* socket address */
 sa_family_t sin_family; /* address family: AF_INET */
 u_int16_t sin_port; /* port in network byte order */
 struct in_addr sin_addr; /* internet address */
};`
- `struct in_addr { /* Internet address */
 u_int32_t s_addr; /* 32 bit address in network byte order */
};`
- **Converting formats: network \leftrightarrow host**
 - unsigned long int **htonl** (unsigned long int hostlong);
 - unsigned short int **htons** (unsigned short int hostshort);
 - unsigned long int **ntohl** (unsigned long int netlong);
 - unsigned short int **ntohs** (unsigned short int netshort);
- **Converting addresses: network \leftrightarrow dotted decimal**
 - long **inet_aton**(char *, struct in_addr *)
 - char ***inet_ntoa**(struct in_addr)

Example of defining a socket address

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>

...
#define MYPORT 22222

...
int s;
struct sockaddr_in sad_loc;
...

sad_loc.sin_family = AF_INET;
sad_loc.sin_port = htons(MYPORT);
inet_aton("127.0.0.1", &sad_loc.sin_addr);
```

Wait for, and accept, connection requests

- **int listen** (int s, int backlog)
 - **s**: local identifier of the socket that receives the connection requests
 - **backlog**: connection requests queue capacity
 - returns immediatly with 0 → OK or -1 → error
- **int accept** (int s, struct sockaddr *addr, socklen_t *addrlen)
 - **s**: local identifier of the socket that waits for the connection requests
 - **addr**: pointer to a *struct sockaddr* with the address of the remote socket
 - **addrlen**: pointer to an integer with the size of *addr
 - Returns an **identifier** of the data communication socket

Used in the server side, only

Requesting connections

- `int connect (int sock, struct sockaddr *s_addr, socklen_t addrlen)`
 - **sock**: local identifier of the client socket
 - **s_addr**: pointer to a *struct sockaddr* with address of the server socket
 - **addrlen**: integer with size of **s_addr*
 - Blocks until the connection is set up (accepted by the server)

Used in the client side, only

Sending and receiving data

- `int send (int s, const void *buf, size_t len, int flags)`
- `int recv (int s, const void *buf, size_t len, int flags)`
 - **s**: identifier of the socket for data communication
 - **buf**: pointer to a buffer
with the data to transmit / where to place the received data
 - **len**: length in bytes of the data to transmit / maximum data to receive
 - **flags**: integer (bit array) specifying several options - just use 0
 - returns number of bytes effectively sent /receive or -1 if error

Used in both sides, client and server

Closing a connection

- **int close (int s)**
 - **s**: local identifier of the socket
 - The socket is destroyed as soon as there is no data in transmission
- **int shutdown (int s, int how)**
 - **s**: local identifier of the socket
 - **how**: 0 - inhibits reception, 1 - inhibits transmission, 2 - inhibits both

Used in both sides, client and server

A few details

- bind, accept, connect - require a socket address of type *sockaddr*
- IPv4 addresses are better specified with type *sockaddr_in*
 - which can be cast to *sockaddr*
- Thus (example with *connect*):

```
struct sockaddr_in serv;  
socklen_t  addrlen = sizeof(serv);  
...  
connect (sock, (struct sockaddr *) &serv, addrlen )
```

A few details

- `send`, `recv` - (TCP) operate on a **stream** model
 - there is no concept of "message".
- Thus, the parts of data that we get with `recv` can be different than those we sent with `send`
 - Ex. `send("abcdefgh")` and in the first `recv` receive "abcd", only, needing another `recv` to get the remainder "efgh".
 - However, this might be irrelevant - depends on the application
- `accept` - allows knowing who is connecting
 - through the socket address that is passed in the parameters
 - but frequently it is not needed

Example of a sever

```
#include ...
#define .....

main() {
    int so, sd, len;
    struct sockaddr_in loc, rem;
    socklen_t addlen = sizeof(loc);
    char buf[BUF_LEN];

    so = socket(PF_INET, SOCK_STREAM, 0);

    loc.sin_family = AF_INET;
    loc.sin_port = htons(MYPORT);
    inet_aton("127.0.0.1", &loc.sin_addr);

    bind(so, (struct sockaddr *) &loc, addlen);

    listen(so, 10);
    sd = accept(so, (struct sockaddr *) &rem, &addlen);

    len = recv(sd, buf, BUF_LEN, 0);
    ...
}
```

Revisit slides 5 & 6

Example of a client

```
#include ...  
#define .....
```

```
main () {  
    int sock, len;  
    struct sockaddr_in serv;  
    socklen_t addlen = sizeof(serv);  
    char buf[BUF_LEN];  
  
    sock = socket(PF_INET, SOCK_STREAM, 0);  
  
    serv.sin_family = AF_INET;  
    serv.sin_port = htons(MYPORT);  
    inet_aton("127.0.0.1",&serv.sin_addr);  
  
    connect(sock, (struct sockaddr *) &serv, addlen);  
  
    scanf("%s", buf);  
    len=send(sock, buf, strlen(buf)+1, 0);  
    ...  
}
```

Revisit slides 5 & 6